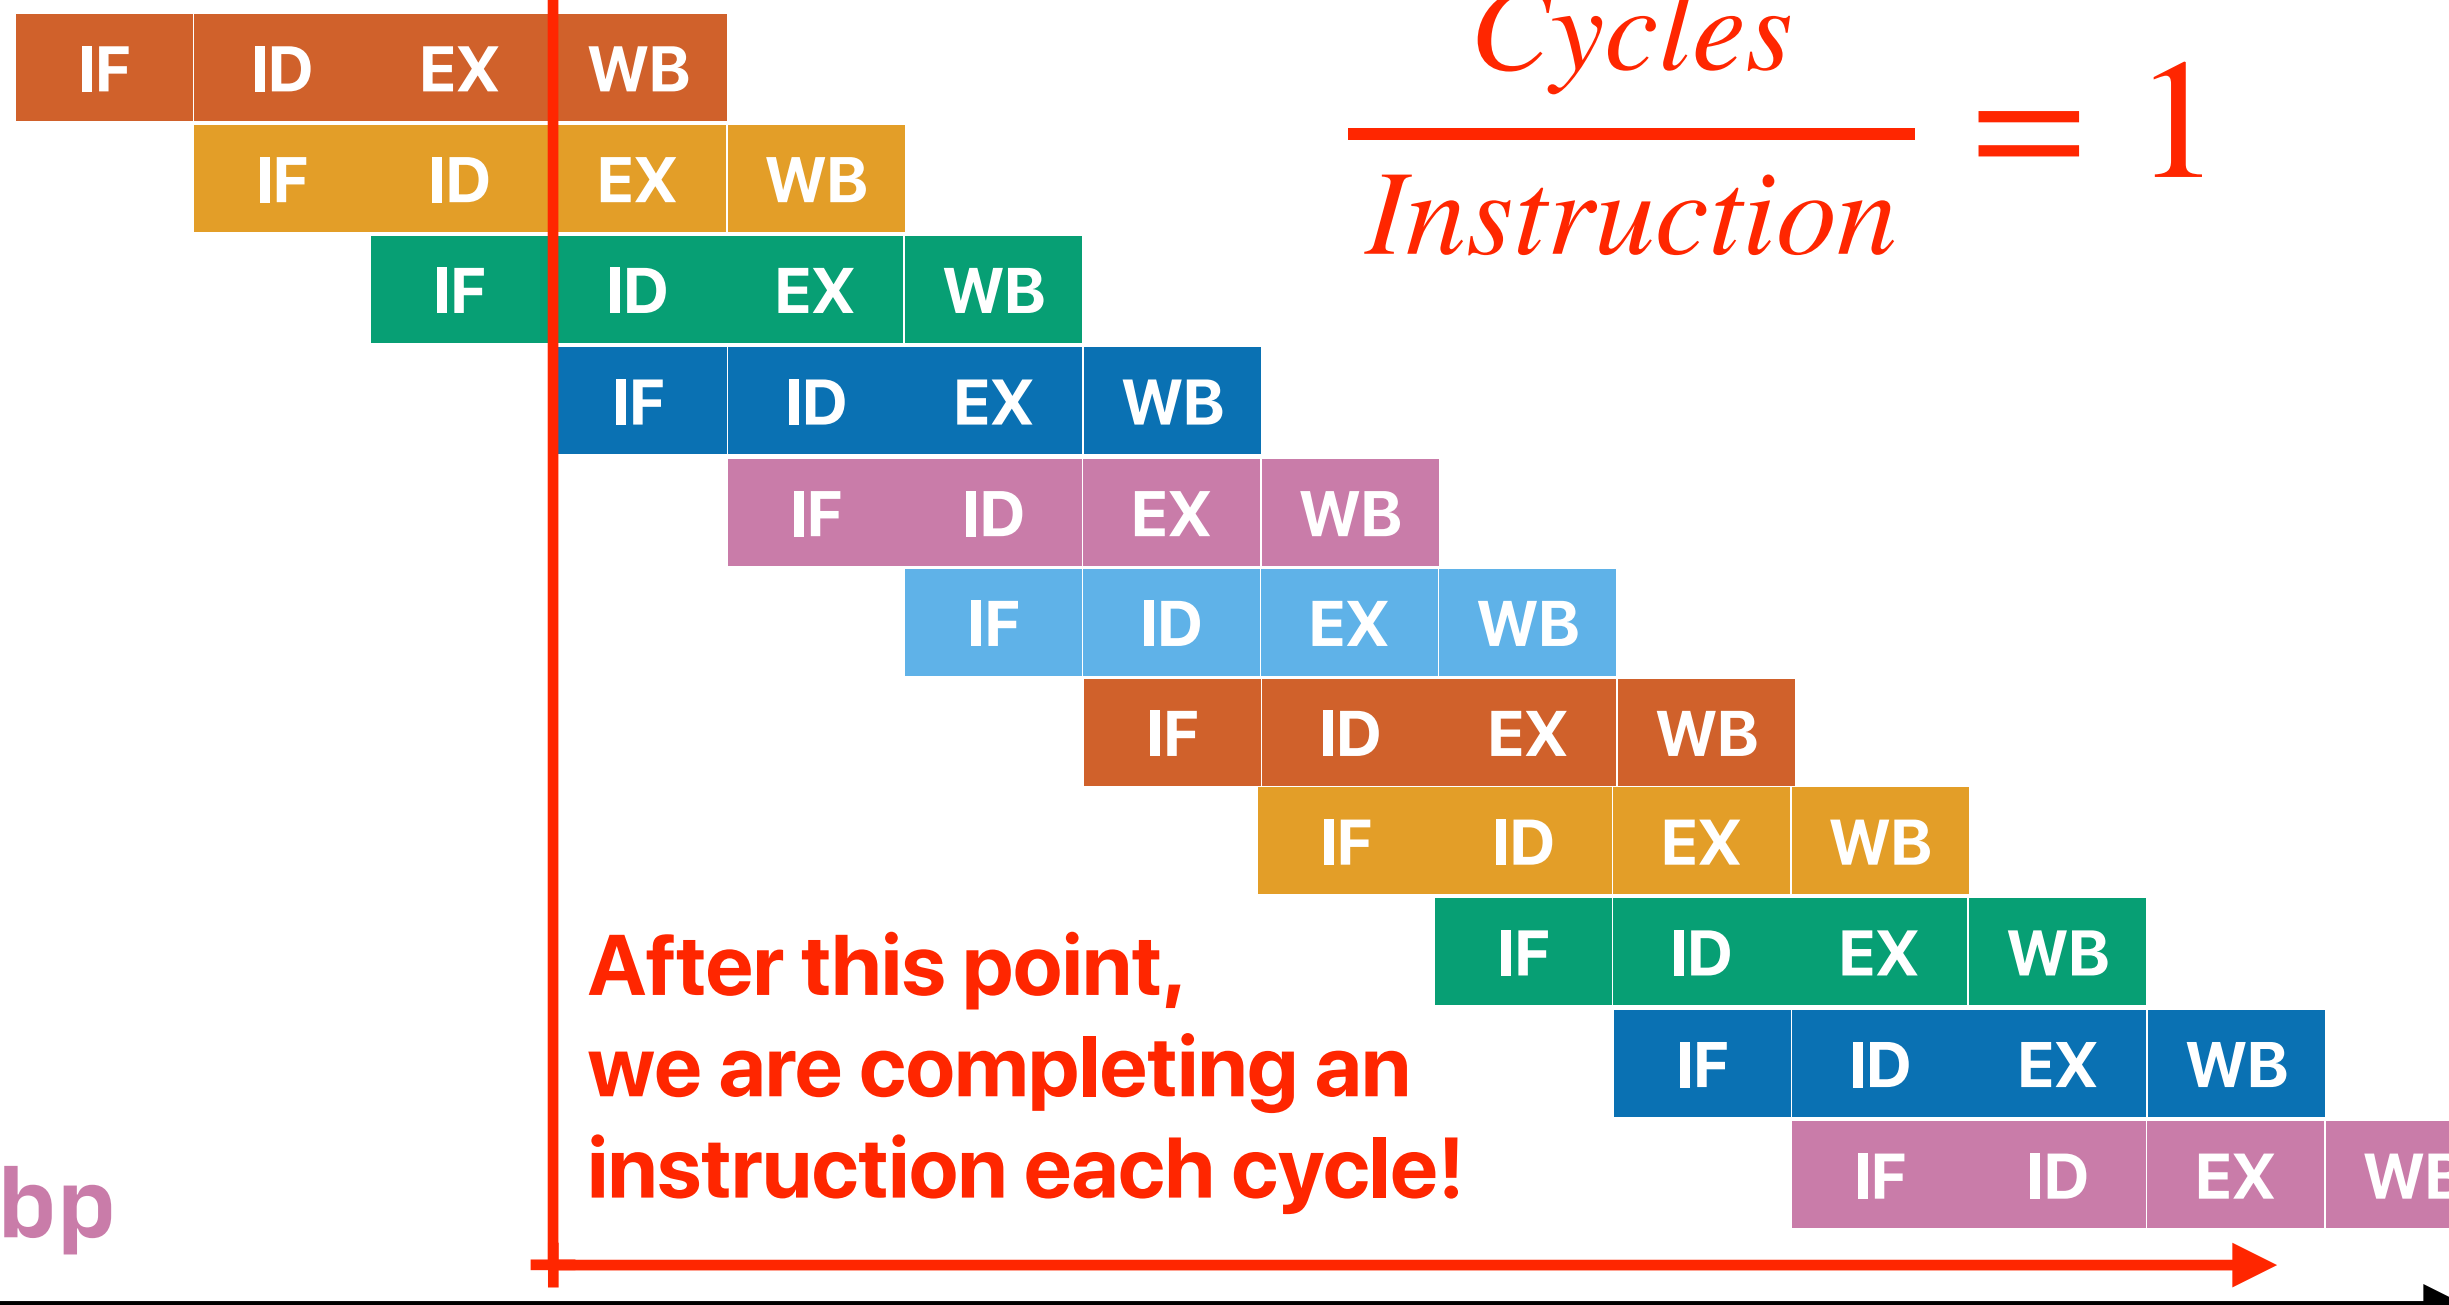


# Modern Processor Design (2): I guess I just feel like...

Hung-Wei Tseng

# Recap: Pipelining

```
addl    %eax, %eax
addl    %rdi, %ecx
addq    $4, %r11
testl   %esi, %esi
movl    $10, %edx
pushq   %r12
pushq   %rbp
pushq   %rbx
subq    $8, %rsp
addl    %rsi, %rdi
movslq  %eax, %rbp
```



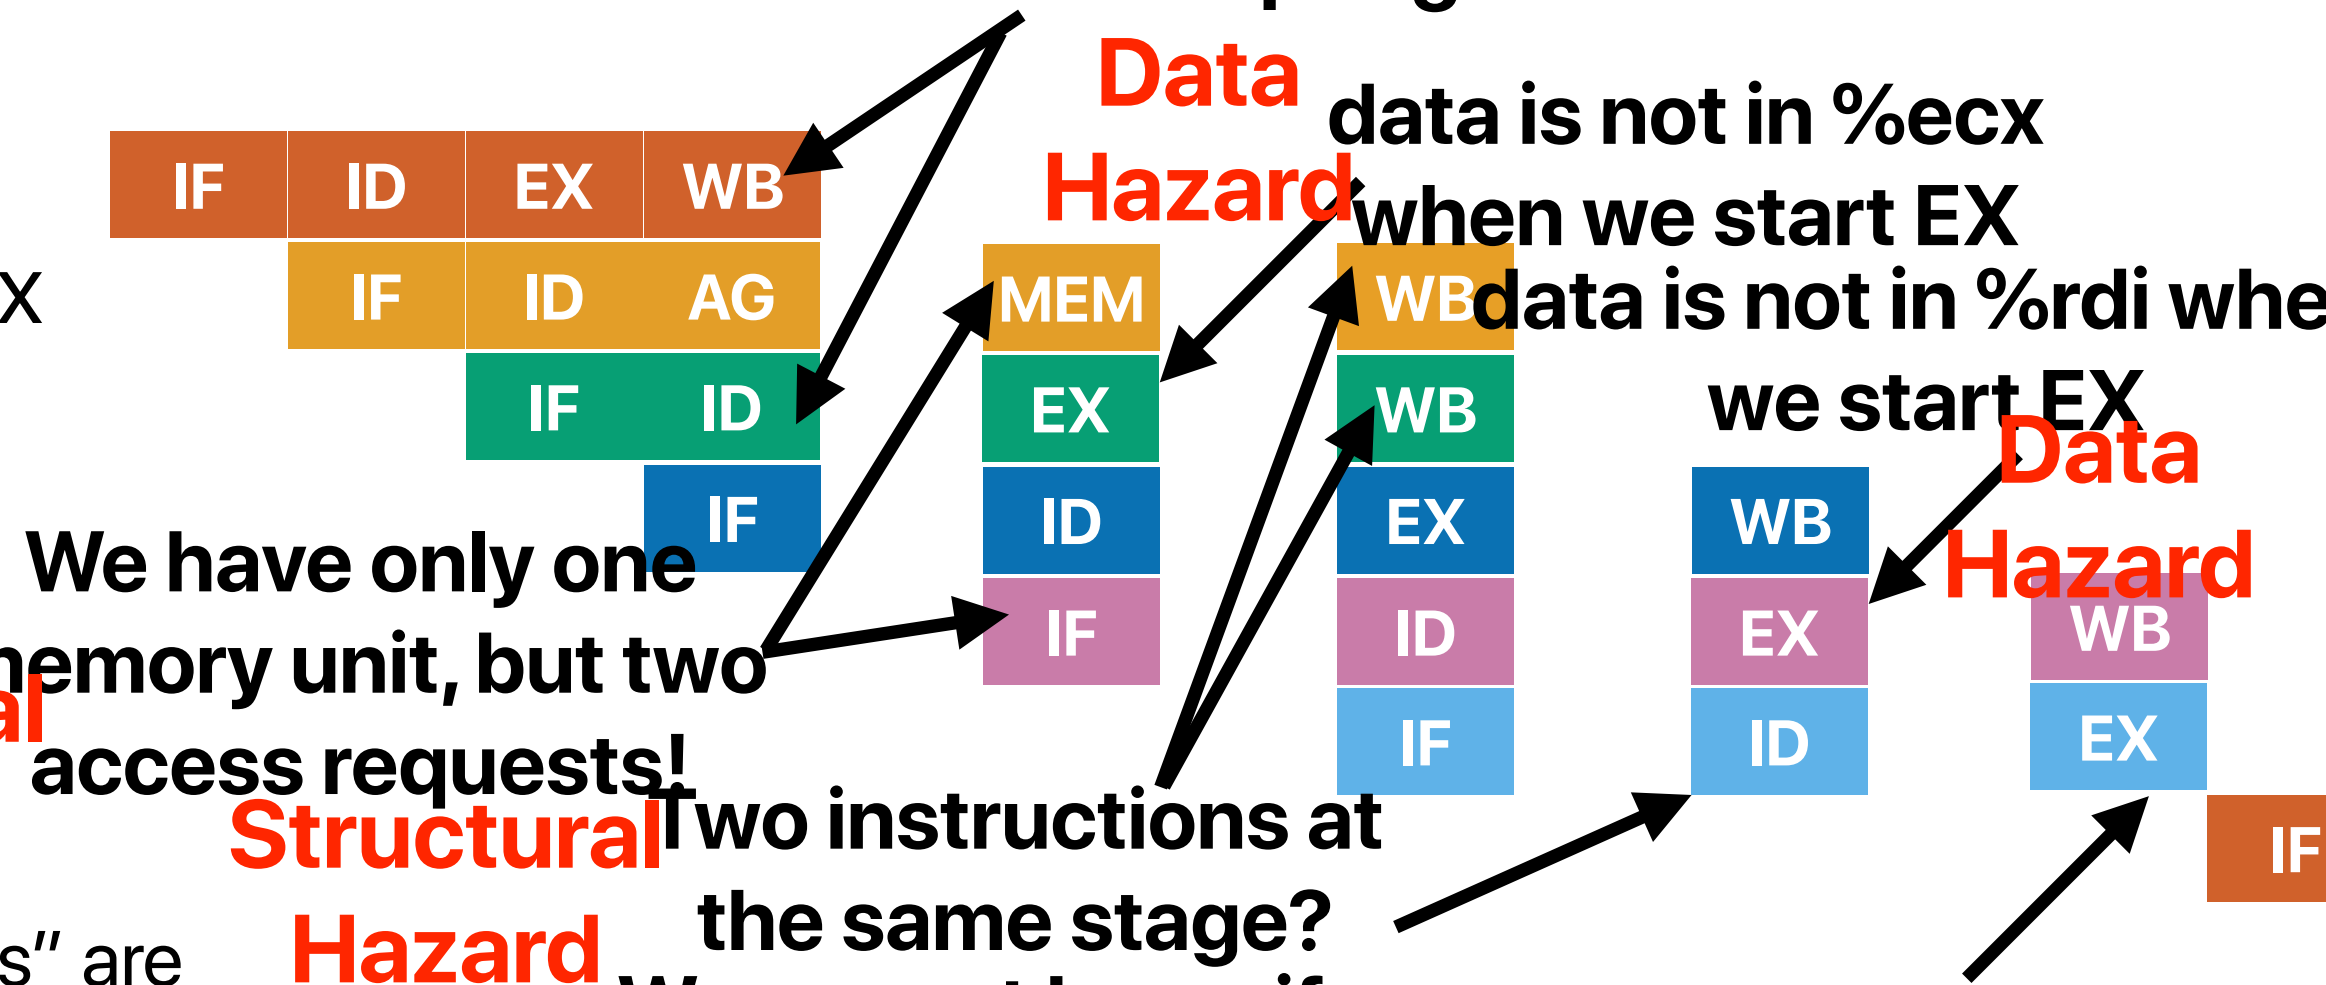
# Pipeline Hazards

```

① xorl %eax, %eax
② movl (%rdi), %ecx
③ addl %ecx, %eax
④ addq $4, %rdi
⑤ cmpq %rdx, %rdi
⑥ jne .L3
⑦ ret
    
```

• How many of the "hazards" are data hazards?

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4



# Recap: Structural Hazards

- Force later instructions to stall
- Improve the pipeline unit design to allow parallel execution
  - Write-first, read later register files
  - Split L1-Cache
  - Non-blocking, multi-banked cache/memory

3:00

**When and how will you make a  
guess?**

# Outline

- Why branch prediction for control hazards
- Dynamic branch predictions
  - Local predictor — 2 bit
  - Global predictor — 2-level
  - Hybrid predictors
    - Tournament
    - Perceptron

# Control Hazards

# How does the code look like?

```
for (unsigned i = 0; i < size; ++i) { // taken when true
```

```
    if (data[i] < threshold) // taken when false
```

```
        call_when_true(&a[i]);
```

```
    else
```

```
        call_when_false(&a[i]);
```

```
}
```

**Branch taken simply means  
we are using branch target  
address as the next address**

.L12:

```
    movl    -32(%rbp), %eax
```

```
    cltq
```

```
    leaq    0(,%rax,8), %rdx
```

```
    movq    -8(%rbp), %rax
```

```
    addq    %rdx, %rax
```

```
    movq    (%rax), %rax
```

```
    andl    $1, %eax
```

```
    testq   %rax, %rax
```

```
    je      .L10
```

```
    movl    -32(%rbp), %eax
```

```
    cltq
```

```
    leaq    0(,%rax,8), %rdx
```

```
    movq    -8(%rbp), %rax
```

```
    addq    %rdx, %rax
```

```
    movq    %rax, %rdi
```

```
    call    call_when_true
```

**Branch taken**

**Branch taken**

```
    jmp     .L11
```

.L10:

```
    movl    -32(%rbp), %eax
```

```
    cltq
```

```
    leaq    0(,%rax,8), %rdx
```

```
    movq    -8(%rbp), %rax
```

```
    addq    %rdx, %rax
```

```
    movq    %rax, %rdi
```

```
    call    call_when_false
```

.L11:

```
    addl    $1, -32(%rbp)
```

.L9:

```
    movl    -32(%rbp), %eax
```


```
    cmpl    -28(%rbp), %eax
```

```
    jl      .L12
```



# Why is "branch" problematic in performance?

① addq \$8, %rbx  
② cmpq %r12, %rbx  
③ jne .L5




The latency of executing the cmpq instruction

We have to wait almost as long as the latency of the previous instruction to make a decision — we cannot fetch anything before that


# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved


# Prediction: What if we guessed right?



# Prediction: What if we are wrong?



# Microprocessor with a "branch predictor"



# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!



# What should branch prediction "predict"

- How many of the following statements are true regarding the why is branch can lead to serious performance issues
  - ① The result value of the previous instruction generating the input to the branch
  - ② The direction of the branch (i.e., taken or not-taken)
  - ③ The target address of the branch
  - ④ The forth-through address of the branch

A. 0  
B. 1  
C. 2  
D. 3  
E. 4

A screenshot of a poll interface. It shows five empty input boxes, each preceded by a letter (A, B, C, D, E) in a small font. The boxes are arranged vertically. The interface has a light blue header and a small '10/1' indicator in the top right corner.

# What should branch prediction "predict"

- How many of the following statements are true regarding the why is branch can lead to serious performance issues

- ① The result value of the previous instruction generating the input to the branch
- ✓ ② The direction of the branch (i.e., taken or not-taken)
- ✓ ③ The target address of the branch
- ④ The forth-through address of the branch

A. 0

B. 1

C. 2

D. 3

E. 4

**What are the "outcome" of the branch?**

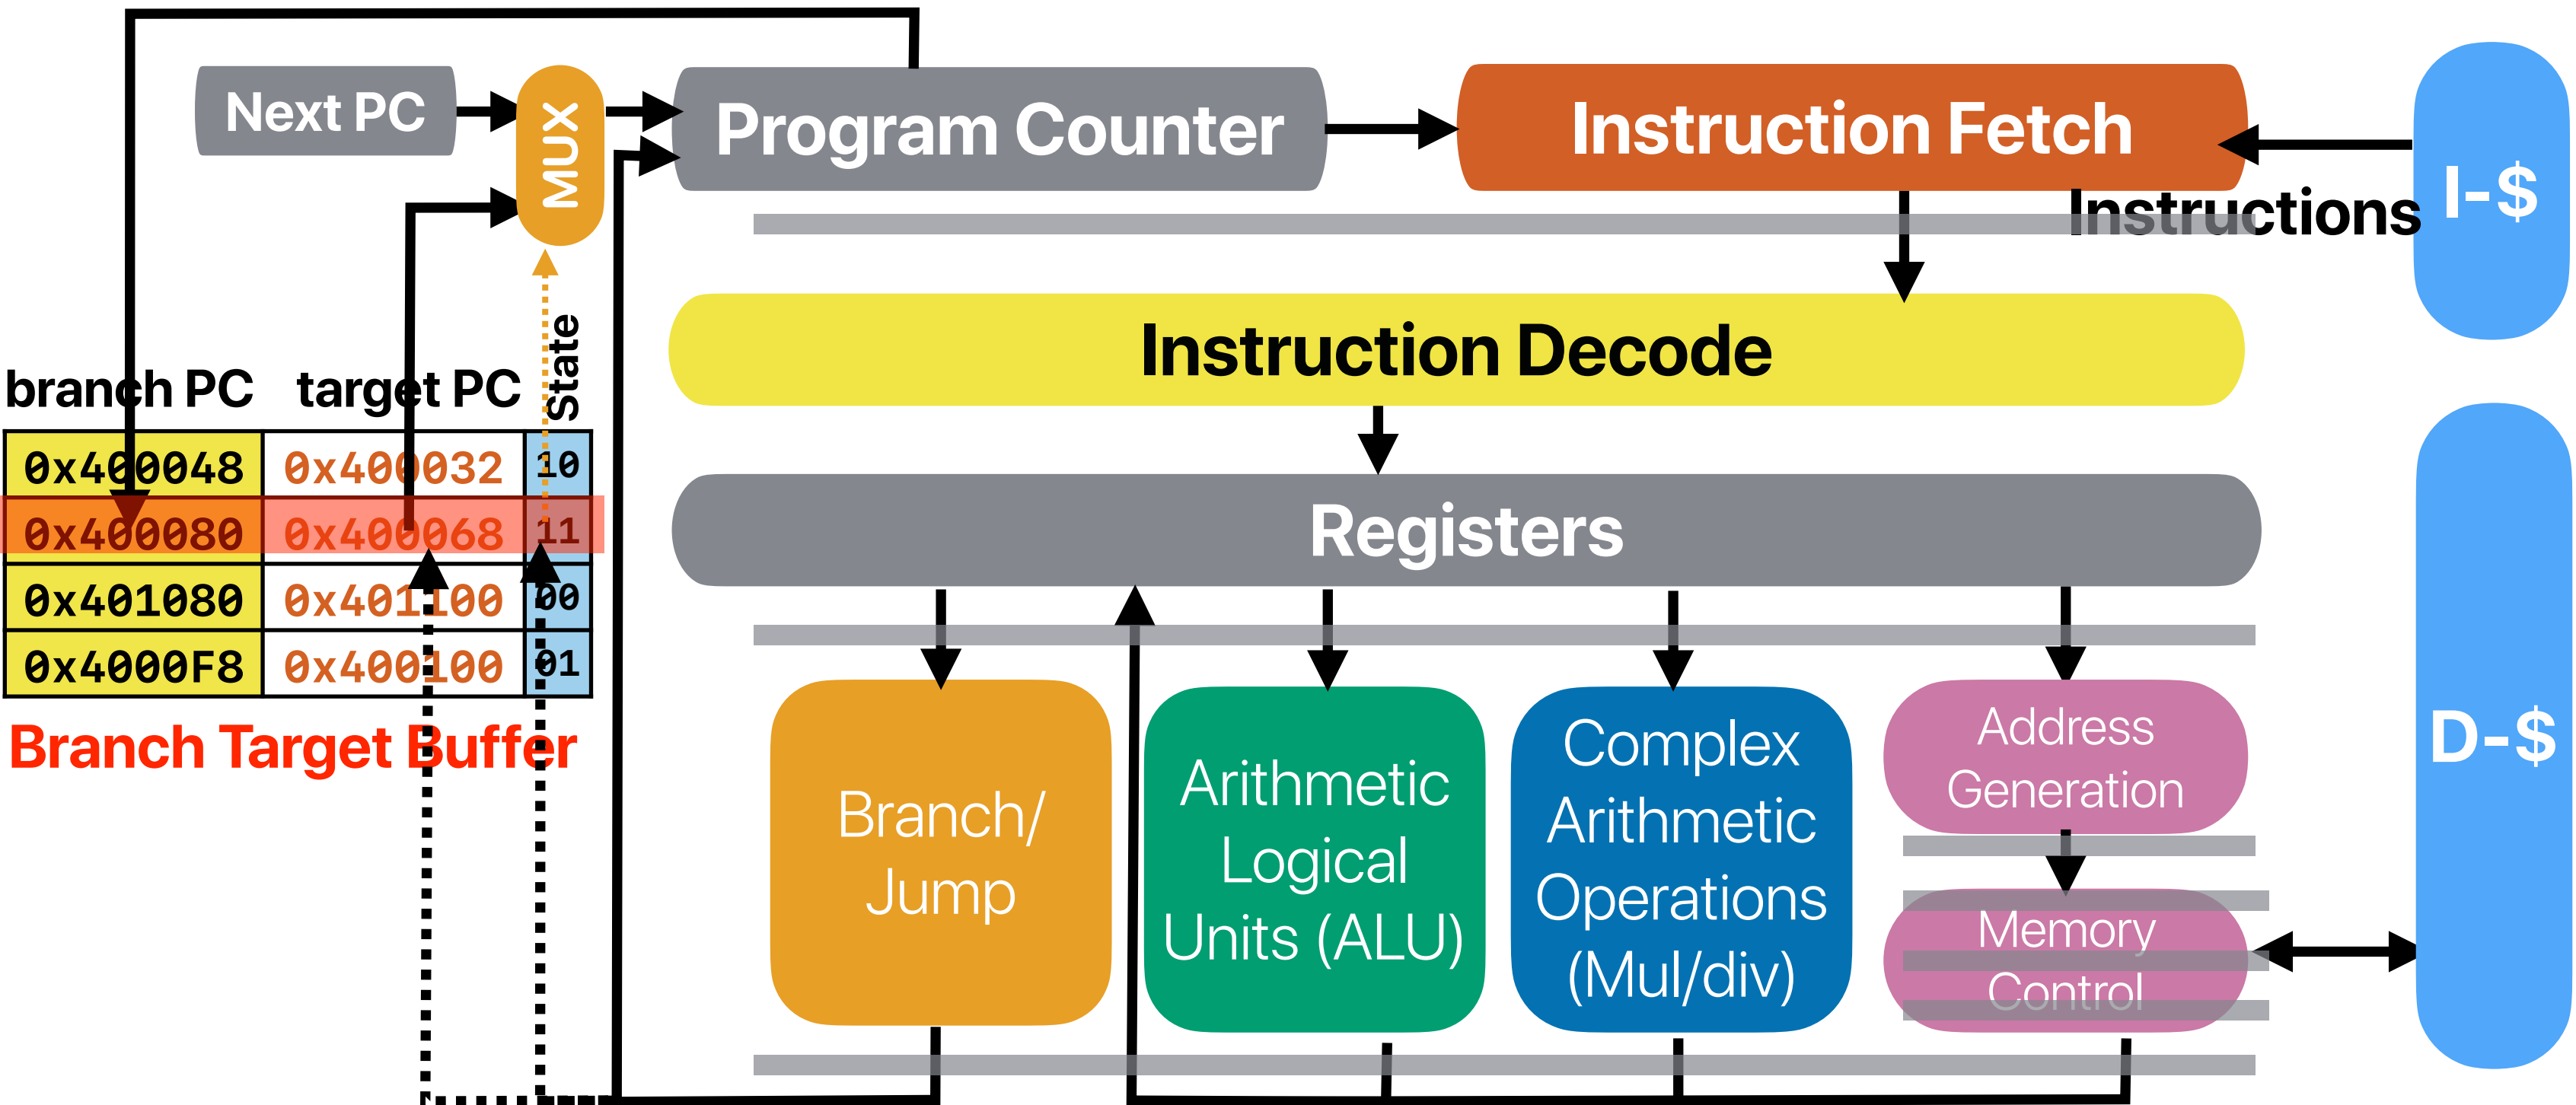
- **Taken, not-take** You need to predict that — history/states

- **Target address, if taken**

**You need a cheatsheet for that — branch target buffer**



# Detail of a basic dynamic branch predictor

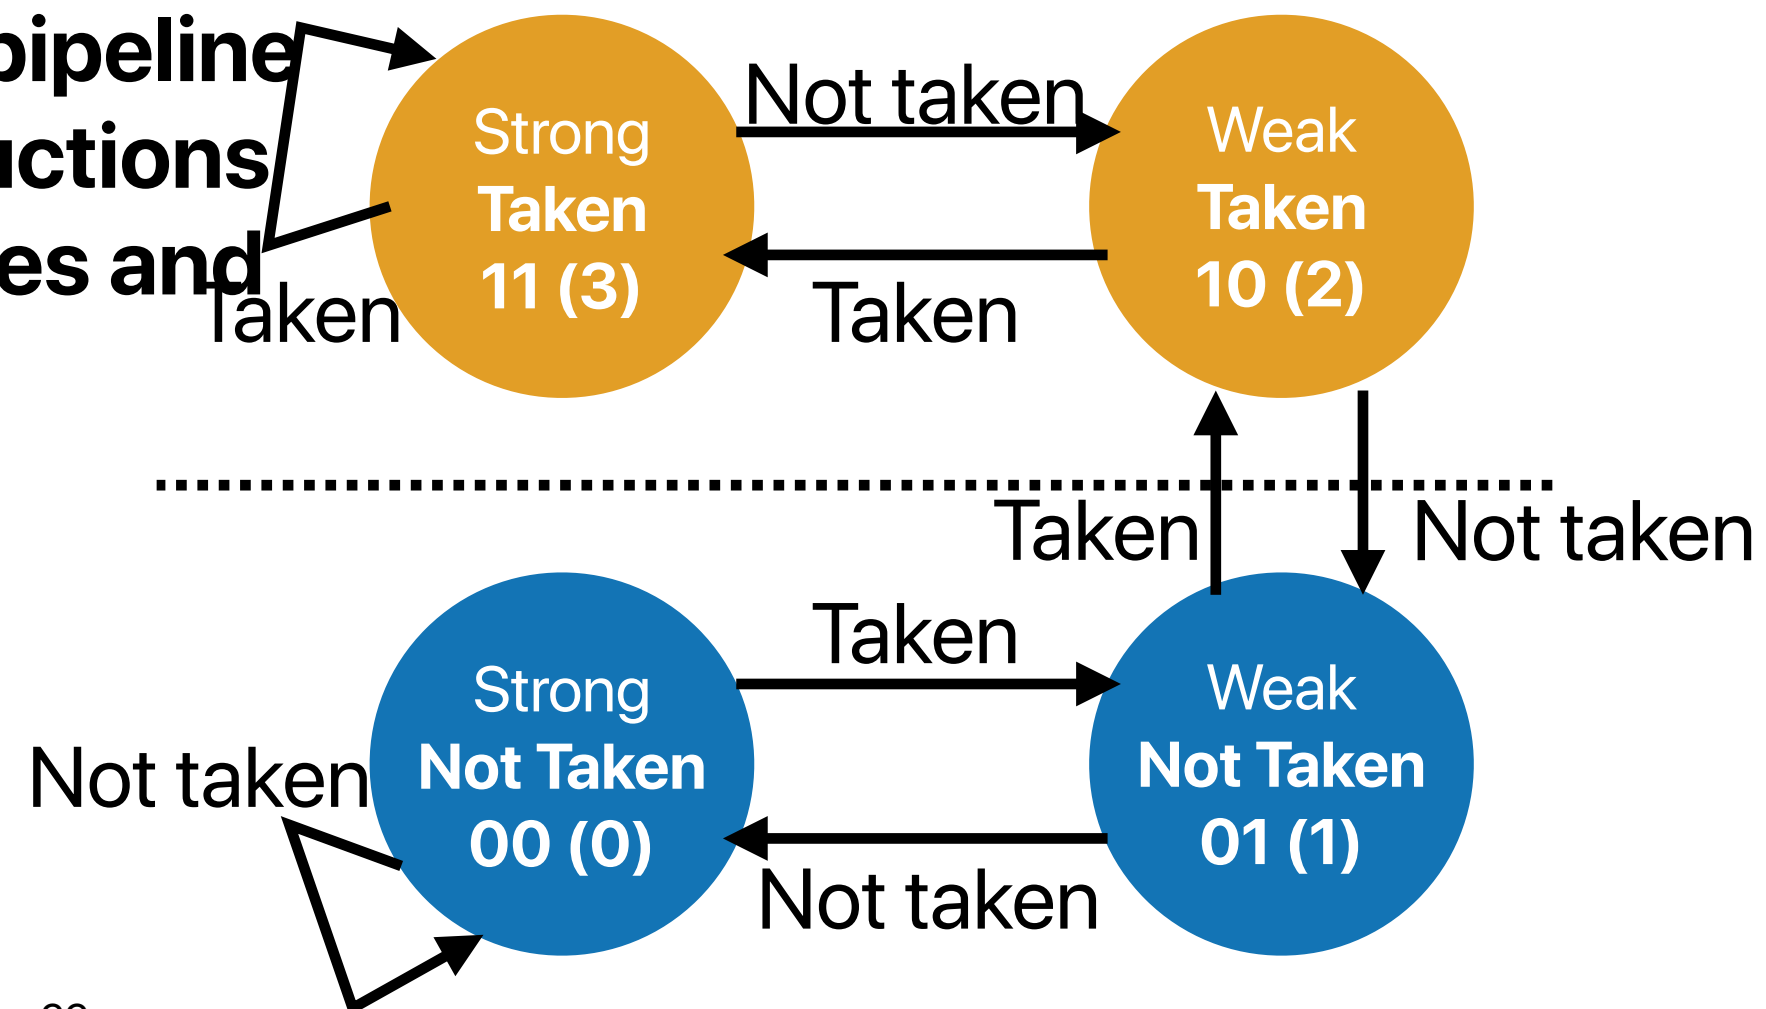


# 2-bit/Bimodal local predictor

- Local predictor — every branch instruction has its own state
- 2-bit — each state is described using 2 bits
- Change the state based on **actual** outcome
- If we guess right — no penalty
- **If we guess wrong — flush (clear pipeline registers) for mis-predicted instructions that are currently in IF and ID stages and reset the PC**

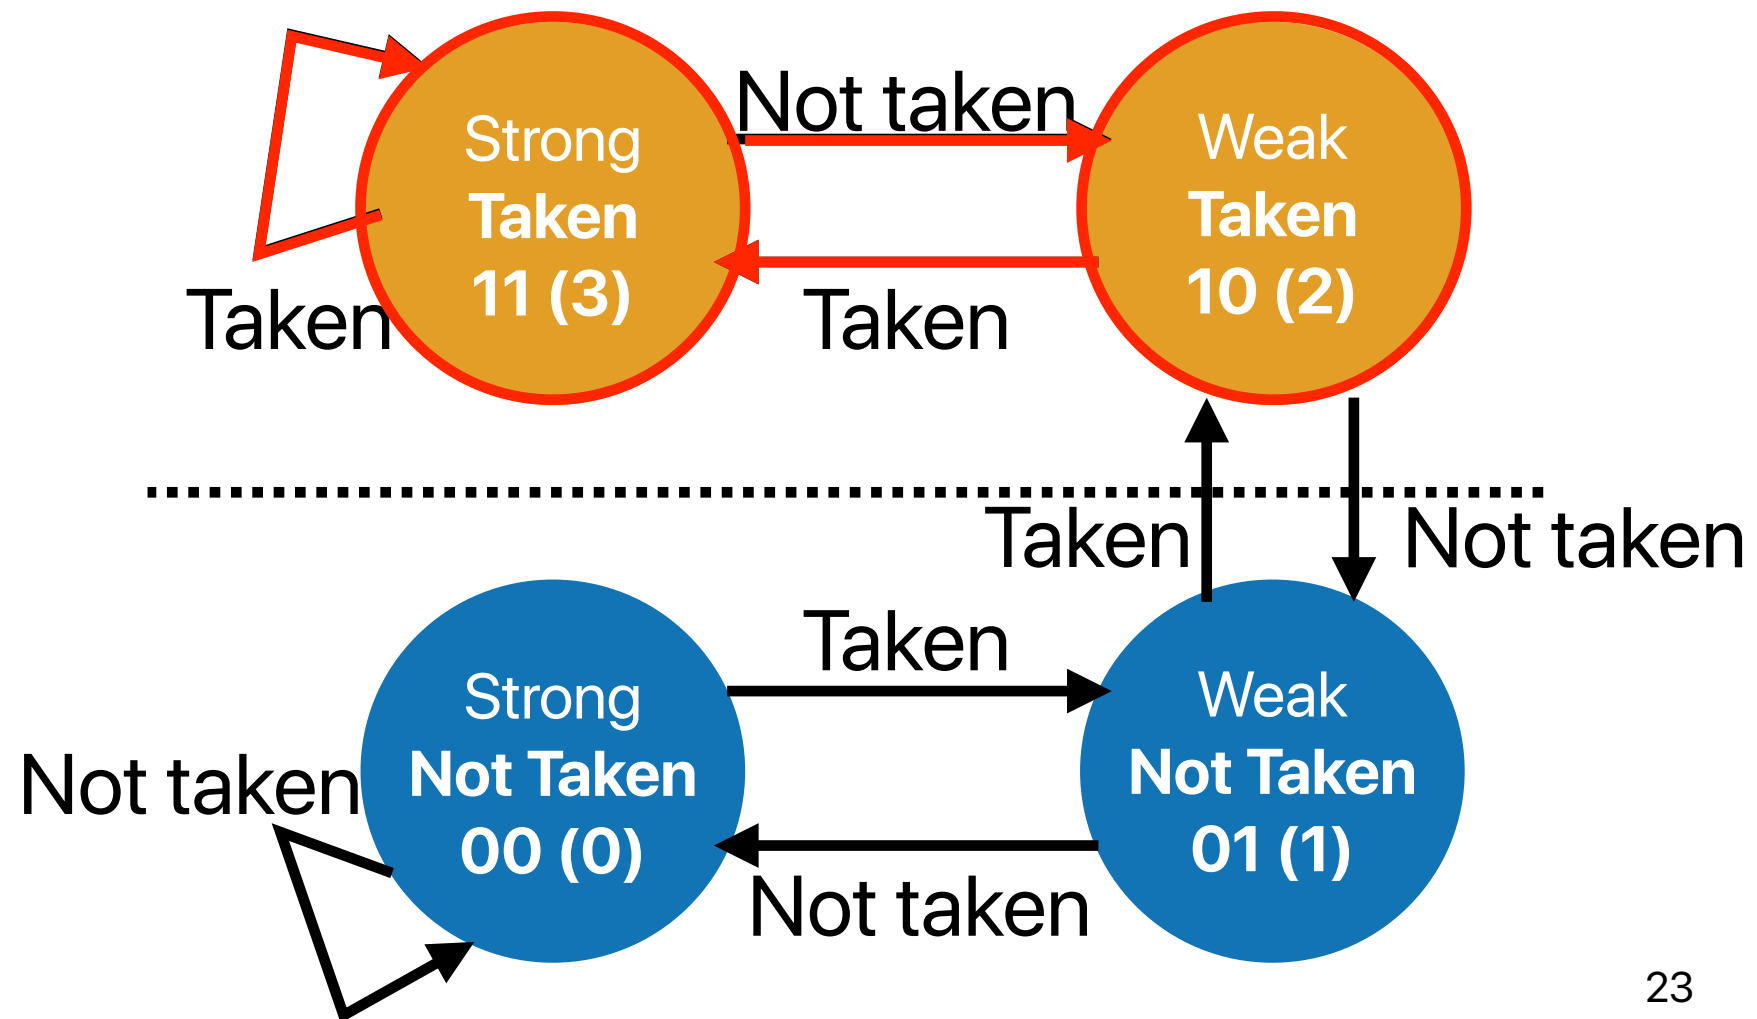
| branch PC | target PC | State |
|-----------|-----------|-------|
| 0x400048  | 0x400032  | 10    |
| 0x400080  | 0x400068  | 11    |
| 0x401080  | 0x401100  | 00    |
| 0x4000F8  | 0x400100  | 01    |

Predict Taken



# 2-bit local predictor

```
i = 0;  
do {  
    sum += a[i];  
} while(++i < 10);
```



| i   | state | predict | actual |
|-----|-------|---------|--------|
| 1   | 10    | T       | T      |
| 2   | 11    | T       | T      |
| 3   | 11    | T       | T      |
| 4-9 | 11    | T       | T      |
| 10  | 11    | T       | NT     |

**90% accuracy!**



# Demo revisited

- Assume that we have a 2-bit local predictor and the values in data is randomly distributed in the number space, what's the branch prediction accuracy of branch X when option is "0" and "1". You may also assume the predictors' states start with 0s.

|   | Without sorting | After sorting |
|---|-----------------|---------------|
| A | 100%            | 0%            |
| B | 50%             | 0%            |
| C | 50%             | 50%           |
| D | 50%             | 100%          |
| E | 0%              | 100%          |

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold) // Branch X
            sum ++;
    }
}
```



# Demo revisited

- Assume that we have a 2-bit local predictor and the values in data is randomly distributed in the number space, what's the branch prediction accuracy of branch X when option is "0" and "1". You may also assume the predictors' states start with 0s.

|   | Without sorting | After sorting |
|---|-----------------|---------------|
| A | 100%            | 0%            |
| B | 50%             | 0%            |
| C | 50%             | 50%           |
| D | 50%             | 100%          |
| E | 0%              | 100%          |

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold) // Branch X
            sum ++;
    }
}
```



# Demo revisited

- Assume that we have a 2-bit local predictor and the values in data is randomly distributed in the number space, what's the branch prediction accuracy of branch X when option is "0" and "1". You may also assume the predictors' states start with 0s.

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold)    // Branch X
            continue;
    }
}
```

|   | Without sorting | After sorting |
|---|-----------------|---------------|
| A | 100%            | 0%            |
| B | 50%             | 0%            |
| C | 50%             | 50%           |
| D | 50%             | 100%          |
| E | 0%              | 100%          |

|   | Without sorting | With sorting |
|---|-----------------|--------------|
| The prediction accuracy of X before threshold | 50%             | 100%         |
| The prediction accuracy of X after threshold  | 50%             | 100%         |

# Demo revisited

If there is no branch predictor on the processor, the code w/ sorting will be slower — but every processor has branch predictors now

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold) // Branch X
    }
}
```

|   | Without sorting | After sorting |
|---|-----------------|---------------|
| A | 100%            | 0%            |
| B | 50%             | 0%            |
| C | 50%             | 50%           |
| D | 50%             | 100%          |
| E | 0%              | 100%          |



|   | Without sorting | With sorting |
|---|-----------------|--------------|
| The prediction accuracy of X before threshold | 50%             | 100%         |
| The prediction accuracy of X after threshold  | 50%             | 100%         |



# How can we evaluate the cost of mis-predicted branches?

- Compare the number of mis-predictions
- Calculate the difference of cycles
- We can get the "average CPI" of a mis-prediction!

## **Demo revisited: evaluating the cost of mis-predicted branches**

- Compare the number of mis-predictions
- Calculate the difference of cycles
- We can get the “average CPI” of a mis-prediction!

**34 cycles on Intel Alder Lake**

**23 cycles on AMD Zen 3**

**Could be more expensive than cache misses**



# 2-bit local predictor

- What's the overall branch prediction (include both branches) accuracy for this nested for loop?

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

(assume all states started with 00)

- A. ~25%
- B. ~33%
- C. ~50%
- D. ~67%
- E. ~75%



# 2-bit local predictor

- What's the overall branch prediction (include both branches) accuracy for this nested for loop?

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

Can we do a better job?

(assume all states started with 00)

- A. ~25%
- B. ~33%
- C. ~50%
- D. ~67%
- E. ~75%**

For branch Y, almost 100%,  
For branch X, only 50%

| i | branch? | state | prediction | actual |
|---|---------|-------|------------|--------|
| 0 | X       | 00    | NT         | T      |
| 1 | Y       | 00    | NT         | T      |
| 1 | X       | 01    | NT         | NT     |
| 2 | Y       | 01    | NT         | T      |
| 2 | X       | 00    | NT         | T      |
| 3 | Y       | 10    | T          | T      |
| 3 | X       | 01    | NT         | NT     |
| 4 | Y       | 11    | T          | T      |
| 4 | X       | 00    | NT         | T      |
| 5 | Y       | 11    | T          | T      |
| 5 | X       | 01    | NT         | NT     |
| 6 | Y       | 11    | T          | T      |
| 6 | X       | 00    | NT         | T      |
| 7 | Y       | 11    | T          | T      |

# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved — 34 cycles on modern intel processors and 23 on AMD processors
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!
- Dynamic branch prediction — predict based on prior history
  - Local predictor — make prediction based on the state of each branch instruction

# Two-level global predictor

Marius Evers, Sanjay J. Patel, Robert S. Chappell, and Yale N. Patt. 1998. An analysis of correlation and predictability: what makes two-level branch predictors work. In Proceedings of the 25th annual international symposium on Computer architecture (ISCA '98).

# 2-bit local predictor

- What's the overall branch prediction (include both branches) accuracy for this nested for loop?

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

(assume all states started with 00)

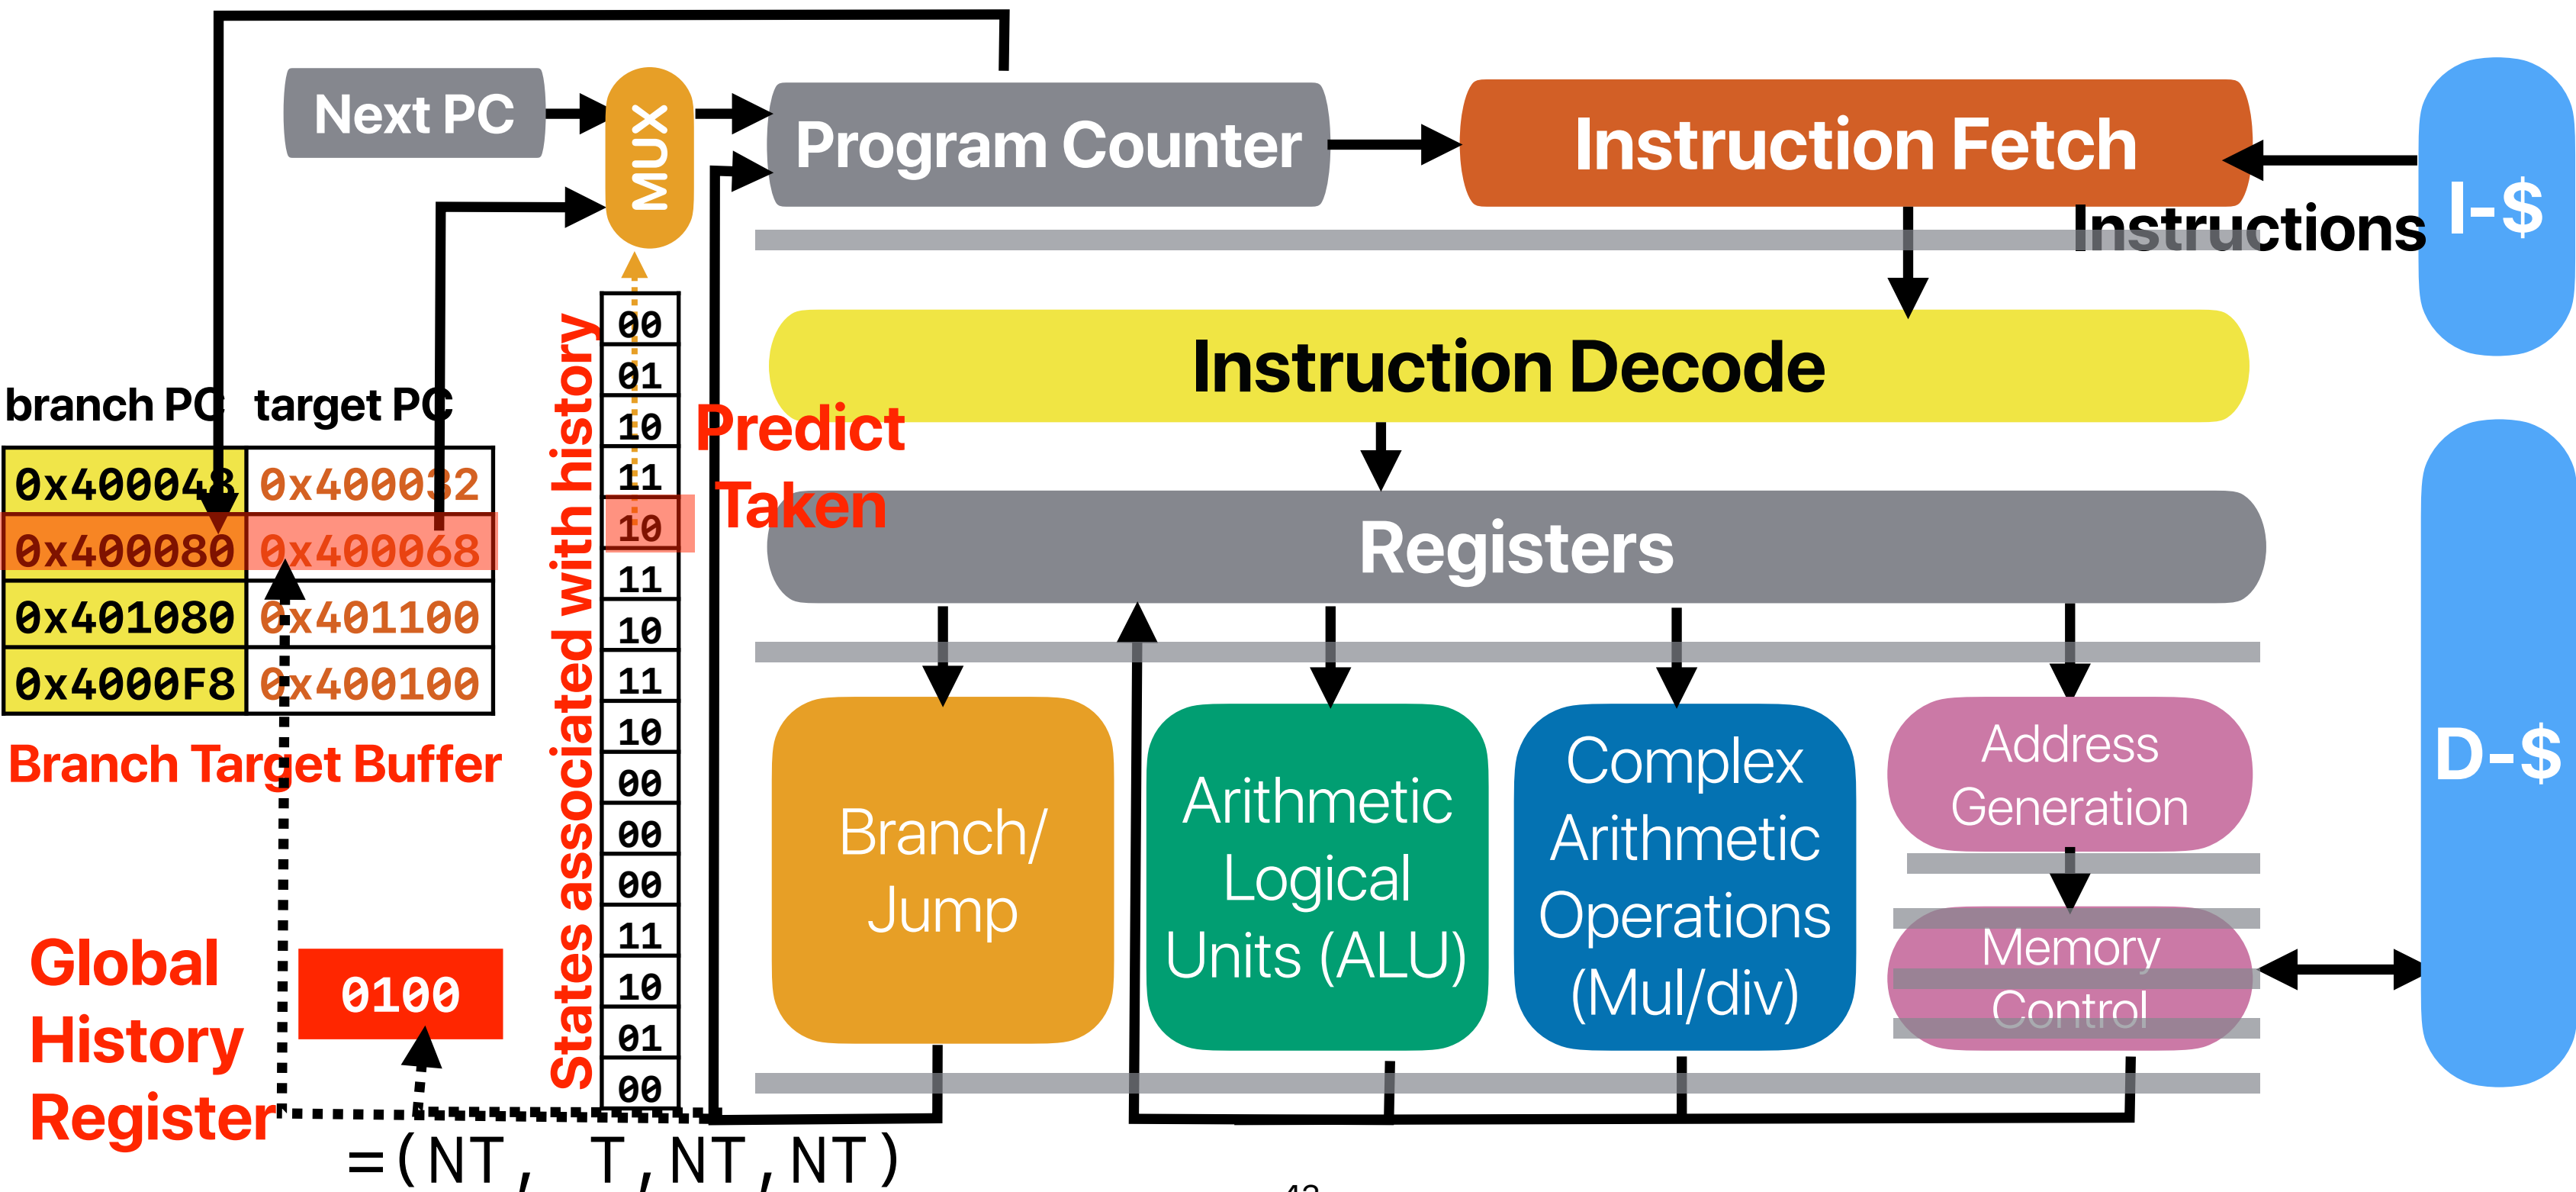
- A. ~25%
- B. ~33%
- C. ~50%
- D. ~67%
- E. ~75%**

**This pattern repeats all the time!**

For branch Y, almost 100%,  
For branch X, only 50%

| i | branch? | state | prediction | actual |
|---|---------|-------|------------|--------|
| 0 | X       | 00    | NT         | T      |
| 1 | Y       | 00    | NT         | T      |
| 2 | X       | 01    | NT         | NT     |
| 2 | Y       | 01    | NT         | T      |
| 2 | X       | 00    | NT         | T      |
| 3 | Y       | 10    | NT         | T      |
| 3 | X       | 01    | NT         | NT     |
| 4 | Y       | 11    | T          | T      |
| 4 | X       | 00    | NT         | T      |
| 5 | Y       | 11    | T          | T      |
| 5 | X       | 01    | NT         | NT     |
| 6 | Y       | 11    | T          | T      |
| 6 | X       | 00    | NT         | T      |
| 7 | Y       | 11    | T          | T      |

# Detail of a basic dynamic branch predictor





# Performance of GH predictor

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

| i  | branch? | GHR | state | prediction | actual |
|----|---------|-----|-------|------------|--------|
| 0  | X       | 000 | 00    | NT         | T      |
| 1  | Y       | 001 | 00    | NT         | T      |
| 1  | X       | 011 | 00    | NT         | NT     |
| 2  | Y       | 110 | 00    | NT         | T      |
| 2  | X       | 101 | 00    | NT         | T      |
| 3  | Y       | 011 | 00    | NT         | T      |
| 3  | X       | 111 | 00    | NT         | NT     |
| 4  | Y       | 110 | 01    | NT         | T      |
| 4  | X       | 101 | 01    | NT         | T      |
| 5  | Y       | 011 | 01    | NT         | T      |
| 5  | X       | 111 | 00    | NT         | NT     |
| 6  | Y       | 110 | 10    | T          | T      |
| 6  | X       | 101 | 10    | T          | T      |
| 7  | Y       | 011 | 10    | T          | T      |
| 7  | X       | 111 | 00    | NT         | NT     |
| 8  | Y       | 110 | 11    | T          | T      |
| 8  | X       | 101 | 11    | T          | T      |
| 9  | Y       | 011 | 11    | T          | T      |
| 9  | X       | 111 | 00    | NT         | NT     |
| 10 | Y       | 110 | 11    | T          | T      |
| 10 | X       | 101 | 11    | T          | T      |
| 11 | Y       | 011 | 11    | T          | T      |

Near perfect after this





# Better predictor?

- Consider two predictors — (L) 2-bit local predictor with unlimited BTB entries and (G) 4-bit global history with 2-bit predictors. How many of the following code snippet would allow (G) to outperform (L)?

—

```
i = 0;
do {
    if( i % 10 != 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100);
```

=

```
i = 0;
do {
    a[i] += i;
} while ( ++i < 100);
```

≡

```
i = 0;
do {
    j = 0;
    do {
        sum += A[i*2+j];
    }
    while( ++j < 2);
} while ( ++i < 100);
```

≥

```
i = 0;
do {
    if( rand() %2 == 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4



# Better predictor?

- Consider two predictors — (L) 2-bit local predictor with unlimited BTB entries and (G) 4-bit global history with 2-bit predictors. How many of the following code snippet would allow (G) to outperform (L)?

about the same

`i = 0;  
do {  
 if( i % 10 != 0)  
 a[i] *= 2;  
 a[i] += i;  
} while ( ++i < 100);`

about the same

`i = 0;  
do {  
 a[i] += i;  
} while ( ++i < 100);`

≡

`i = 0;  
do {  
 j = 0;  
 do {  
 sum += A[i*2+j];  
 }  
 while( ++j < 2);  
} while ( ++i < 100);`

L could be better

≥

`i = 0;  
do {  
 if( rand() %2 == 0)  
 a[i] *= 2;  
 a[i] += i;  
} while ( ++i < 100)`

A. 0

B. 1

C. 2

D. 3

E. 4

# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved — 34 cycles on modern intel processors and 23 on AMD processors
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!
- Dynamic branch prediction — predict based on prior history
  - Local predictor — make predictions based on the state of each branch instruction
  - Global predictor — make predictions based on the state from all branches
  - Both are not perfect

# Announcements

- Reading quiz #7 due next Tuesday
- Programming assignment #2 and assignment #3 due this Thursday

# Computer Science & Engineering

# 203

# つづく

